

Terrain in Battlefield 3: A modern, complete and scalable system

Mattias Widmark

Software Engineer, EA Digital Illusions (DICE), Stockholm

Overview

- Scalability – hierarchies, payloads and limitations
- Workflows – realtime in-game editing
- CPU and GPU performance
- Procedural virtual texturing – powerful GPU optimization
- Data streaming – minimizing memory footprint
- Robustness – global prioritization
- Procedural mesh generation
- Conclusions

But first...

- ... what are we talking about?
- Frostbite terrain has many aspects other than the terrain mesh itself
 - Let's look at them!

- 
- The background image shows a vast desert landscape with rugged, light-colored mountains under a clear blue sky. In the foreground, there is a cluster of large, cylindrical industrial storage tanks, possibly for oil or water, situated near a winding road. The terrain is arid with sparse, dry vegetation. A dark rectangular box containing the text is positioned in the lower-left area of the image.
- Heightfield-based
 - Mesh procedurally generated at runtime



- Surface rendering with *procedural shader splatting*
 - Arbitrary shaders splatted according to painted masks



- Spline and quad decals

- *Terrain decoration*
 - Automatic distribution of meshes (trees, rocks, grass) according to mask
 - Billboards supported



- Terrain decoration
 - Important as the terrain surface itself





- Destruction/dynamic terrain
 - Destruction depth map
 - Controls crater depth around ie static models
 - Physics material map
 - Controls surface effects, audio, crater depth and width



- Rivers/lakes
 - Implemented as free-floating decals
 - Water depth in pixel shader

Terrain raster resources


- Multiple raster resources used
 - Heightfield
 - Shader splatting masks
 - Colormap, used as an overlay on top of shader splatting
 - Physics materials
 - Destruction depth mask
 - Albedo map for bounce light
 - Additional mask channels

Overview

- **Scalability – hierarchies, payloads and limitations**
- Workflows – realtime in-game editing
- CPU and GPU performance
- Procedural virtual texturing – powerful GPU optimization
- Data streaming – minimizing memory footprint
- Robustness – global prioritization
- Procedural mesh generation
- Conclusions

Scalability


- Our definition of *scalability*
 - Arbitrary view distance (0.06m to 30 000m)
 - Arbitrary level of detail (0.0001m and lower)
 - Arbitrary velocity (supercars and jets)
- Main observation
 - It is all about hierarchies!
 - Consistent use of hierarchies gives scalability “for free”
- Hierarchies not new to terrain rendering
 - Frostbite approach similar to flight simulators
- *Quadtrees* hierarchies used for all spatial representations
- Assuming knowledge of quadtrees, we jump right into Frostbite specifics!



Quadtree node payload

- The *node payload* is a central concept
- A quadtree node can be attributed with a "data blob"; the payload
- Payload is
 - a tile of raster data
 - a cell of terrain decoration
 - A list of instances (rock, grass, trees)
 - a piece of decal mesh
- All nodes have payload...
 - ... but only a few have it loaded

Payload (red dot)




Nodes with and without payload

- Payloads are constantly in motion
 - They are loaded (streamed), generated or freed every frame
 - Only a fraction of the nodes have payload resident
- Payload movement is governed by prioritization mechanisms...
 - ... but more of that later

LOD payload

- Non-leaf nodes have payload too
- These payloads are used as LODs
- Detail level depends on payload depth
 - Nodes closer to root represent lower detail

Payload (red dot)
LOD payload (orange dot)




Increasing LOD

View-dependent payload usage


Set of payloads (green dots) used for a certain observer position

- Note area to the left is distant and use lower LOD




Observer moves and another set is used

- Area to the left now use higher LOD




Motivation for LOD payloads

- With LOD
 - cost (payload count) is mostly independent of terrain size
 - Scalable!
- Without LOD
 - cost depend on terrain size
 - Not scalable!



Generation of raster LOD payload

- Source data and workflows on leaf level
- LODs generated automatically by pipeline
- Requirements
 - Tile overlap (borders) for rendering algorithms
 - Continuity
- Recursive (reverse) depth first algorithm
 - Green LOD tile is generated
 - Four children (red) and up to 12 neighbor tiles (blue) are used



Terrain decoration payloads

- Terrain decoration payload
 - Is a list of instance transforms (for grass, trees, rocks)
 - Is generated at runtime according to
 - scattering rules
 - shader splatted masks (position/density)
 - Note that we allow shaders to modify masks!
 - heightfield (ground-clamping and orientation)

Terrain decoration LOD payloads

- Quite unique (and slightly confusing) concept
- LOD payloads used for scalability
 - Near-root payloads provide high view distance
 - Near-leaf payloads provide high density
- Payloads can overlap
 - Providing high view distance *and* density



Overview

- Scalability – hierarchies, payloads and limitations
- **Workflows – realtime in-game editing**
- CPU and GPU performance
- Procedural virtual texturing – powerful GPU optimization
- Data streaming – minimizing memory footprint
- Robustness – global prioritization
- Procedural mesh generation
- Conclusions

A screenshot of the FrostEd terrain editor interface. The main view shows a 3D rendering of a park-like environment with a paved path, trees, and a body of water. A red arrow points to the 'Game View' button in the top-left corner of the central window. The left panel contains a terrain browser with categories like 'Terrain', 'Destruction depth', 'Color', 'Physics', 'Heightfield', 'Terrain layers', and 'Terrain layouts'. The 'Terrain layers' section lists various materials such as 'MP_Subway_Mud', 'Backdrop Trees', 'Underground Rubble', 'Gravel', 'Flowerbed White' (which is selected), 'Flowerbed Pink', 'Dirt (disabled)', 'Grass Dry (disabled)', 'Grass Tall with Rubbish', 'Underwater', and 'Grass'. The right panel includes 'Tool settings' for brush parameters like position, height, magnification, and pressure, as well as a 'Raster brushes' palette with several circular brush icons. The bottom right panel displays 'Properties' for the selected terrain layer, including scattering, shading, and debugging information.

External tools

- When tools are not enough terrain can be exported and imported
 - Select all or part of terrain
 - Metadata + raw file
 - Edit raw file and reimport
 - Metadata will import to right area
 - Puts *WorldMachine*, *GeoControl* in the loop
 - A common workflow

Workflow issues


- Conflict between data compression and realtime editing
- Realtime editing bypass pipeline
- Clever update scheme for procedural content needed
 - We already had one (destruction)
- Frostbite terrains too large for some popular
 - GeoControl and WorldMachine do not like 8k+



Overview

- Scalability – hierarchies, payloads and limitations
- Workflows – realtime in-game editing
- **CPU and GPU performance**
- Procedural virtual texturing – powerful GPU optimization
- Data streaming – minimizing memory footprint
- Robustness – global prioritization
- Procedural mesh generation
- Conclusions

Efficient on CPU



- All work done in jobs, most on SPU and many wide
 - Early unoptimized versions consumed 10ms+ PPU time
 - BF3 final measurements (PS3)
 - 1-2ms SPU (peaks at ~8ms when lots of terrain decoration is happening)
 - <1ms PPU

Efficient on GPU

- Pre-baked shader permutations to avoid multi-pass
- *Procedural virtual texturing* exploit frame-to-frame coherency
- Typical figures (PS3):
 - Full screen GBuffer laydown (w/o *detail overlay* and terrain decoration): 2.5-3ms
 - Full screen GBuffer laydown (w/ detail overlay): 2.5-7ms
 - Terrain decoration: 1-4ms
 - Virtual texture tile compositing: 0.2-0.5ms

Overview

- Scalability – hierarchies, payloads and limitations
- Workflows – realtime in-game editing
- CPU and GPU performance
- **Procedural virtual texturing – powerful GPU optimization**
- Data streaming – minimizing memory footprint
- Robustness – global prioritization
- Procedural mesh generation
- Conclusions

GPU optimization: Procedural virtual texturing

- Motivations
 - With shader splatting, artists can create beautiful terrains...
 - ... rendering very slowly (10-20ms)
 - Shader splatting not scalable in view distance
 - Cannot afford multi-pass
- By splatting into a texture
 - we leverage frame-to-frame coherency (performance)
 - can render in multiple passes (scalability)
- Rendering full screen using the texture cost 2.5-3ms (PS3)

Virtual texture key values

- 32 samples per meter
- 256x256 tiles with integrated two pixel border
- Atlas storage with default size 4k x 2k
- Two DXT5 textures

R	G	B	A	R	G	B	A
Diffuse R	Diffuse G	Diffuse B	Normal X	Smoothness/ Destruction	Normal Y	Specular	Normal Z

- Very large, can easily reach 1M x 1M (= 1Tpixel)!
 - Typical virtual textures are 64k x 64k


Indirection texture format

- RGBA8
- Indices into virtual texture tile atlas
- Scale factor for low-res areas...
 - ... where a tile covers more than one indirection sample
- CLOD fade factor
 - Used to smoothly fade in a newly composited tile (*fade-to* tile)
 - Previous LOD (*fade-from* tile) is already in atlas and fetched using indirection mips
 - CLOD factor updated each frame

R	G	B	A
Index X	Index Y	Scale	CLOD fade

The “Teratexture”

- How do we reach 1M x 1M?
 - Indirection texture can easily go 4k x 4k
 - Way too large!
- *Clipmap* indirection texture
 - Clipmap – early virtual texture implementation
 - Replace 4k indirection texture with 6 64x64 clipmap layers



Clipmap indirection

- Clipmap level is resolved on CPU for each draw call
 - Avoids additional pixel shader logic
 - Requires each 64x64 map to have its own mip chain
- Texture space has to be roughly organized with world space
 - Not an issue for terrain
 - More generic use cases are probably better off with multiple classic virtual textures

Tile compositing


- Tiles are composited on GPU and compressed on GPU or SPU
- Benefits (compared to streaming from disc)
 - Small disc footprint – data magnification
 - Source raster data magnified ~1000x
 - Low latency
 - Tile is ready to use next frame
 - Dynamic updates
 - Destruction
 - Realtime editing
 - Efficient workflow
 - Artists don't have to paint hundreds of square kilometers down to every last pebble

Virtual texture issues

- Blurriness (consoles have too small tile pool)
- Runtime texture compression quality
- GPU tile compositing cost offset some of the gain
 - ~0.25ms/frame on Xenon
- Limited hardware filtering support
 - Expensive software and/or lower-quality filtering
- Compositing latency
- Tile compositing render target memory usage


Virtual texture issues

- Virtual texture has practical limit at around 32 samples per meter
 - *Detail shader splatting* fills in to required sharpness (500-1000 samples per meter)
- We now have two shader splatting methods
 - Diffuse/Normal/Specular/Smoothness splats into virtual texture
 - Direct* splats details into Gbuffer
- Performance concerns
 - Have to limit detail view distance
 - Typically 50m-100m



Virtual texture issues

- Indirection update performance
 - Updating indirection maps is expensive
 - 4-6 64x64 maps with mips
 - Do in job (SPU)
 - Update only dirty areas
 - New tile
 - CLOD fade factor
 - Recentering when clipmap region moves
 - Wraparound
 - Only edges affected on recentering (bottom image)



Overview

- Scalability – hierarchies, payloads and limitations
- Workflows – realtime in-game editing
- CPU and GPU performance
- Procedural virtual texturing – powerful GPU optimization
- **Data streaming – minimizing memory footprint**
- Robustness – global prioritization
- Procedural mesh generation
- Conclusions

Data streaming


- Frostbite 1 did not stream terrain
- Streaming required for larger *Battlefield 3* and *NFS: The Run* levels

Streaming basics

- Streaming unit: Raster tiles (aka node payloads)
 - Typical tile sizes
 - Heightfield: 133x133x2 bytes
 - Mask: 66x66x1 bytes x 0-50 tiles per payload
 - Color: 264x264x0.5 bytes
- Fixed-size tile pools (atlases)
 - Typical atlas sizes
 - Heightfield: 2048x2048
 - Mask: 2048x1024
 - Color: 2048x2048

Streaming modes

- Tile-by-tile (aka *free*) streaming
 - For slower gameplay
- Tile *bundle* (aka *push-based*) streaming
 - For faster games
 - Tiles associated with a layout are bundled
 - Based on *terrain resolution layouts*
- Hybrid streaming (most common)
 - Bundles used at selected spawn points and transitions
 - Free streaming fill in the rest



Layout of all data on level

Data subset loaded at spawn point

A *Terrain resolution layout* defines the subset

Subsets loaded (and unloaded) as user progresses
through level

Resident set size: Powerful *blurriness*


- Streaming does not remove memory footprint
 - A *resident set* is still needed
- Resident set can be huge
 - Theoretical value for BF3 level on PS3: 70+Mb!
- *Blurriness* feature shrinks resident set significantly
 - Increase blurriness by one and save around 70%
 - Very important memory saver for BF3
 - Shipped with 32Mb terrain resident set

Blurriness

- Blurriness is mip bias applied to terrain raster streaming
 - Blurriness = 0:
 - Streaming continues until raster is sharp
 - Blurriness = 1:
 - Streaming stops when raster is slightly blurry (texel covers 2x2 pixels)
 - Blurriness = 2:
 - Streaming stops when raster is significantly blurry (texel covers 4x4 pixels)

Blurriness: Implementation by pipeline trick

- For each level of blurriness
 - cut tile size in half
 - add one leaf level
- No data is lost – it is only shifted downwards




Blurriness 0

256x256-sized tiles

8 **resident tiles** (500kpixel)

13 nodes



Blurriness 1

128x128-sized tiles

14 **nodes added**

10 **resident tiles** (164kpixel, saves 68%)

27 nodes

BF3 blurriness use case

- Blur expensive rasters
 - Heightfield (2 bytes per sample)
 - Mask (1 byte per sample)
- Keep cheap rasters
 - Colormap (DXT1, 0.5byte per sample)
- Put detail (for example occlusion) in colormap to hide blurry heightfield/normal map

FPS: 63.24 (0)
Active soldiers/actors: 1/0 (max 24)
Active vehicles: 6 (max 6)
Total vehicles: 33 (max 12)




- **Physics**
 - Wrong streaming setup gives strange artifacts
 - Vehicles spawning in air or in ground
 - Invisible and disappearing walls and holes

Reducing disc seeks


- Often main reason for latency
 - Can seek maybe four files per second
 - A terrain can have hundreds of files (tiles/payloads)
- Methods to reduce the number of seeks
 - Use terrain resolution layouts at spawn
 - Otherwise **minutes** can pass before stabilization (waiting for file seeks)!
 - Co-locate overlapping tiles of different types
 - Store heightfield tile together with color and mask tiles

Reducing disc seeks

- Methods to reduce the number of seeks (cont'd)
 - Co-locate nearby tiles
 - Group leaf node payloads as *second LOD* in ancestor node
 - Saves 20-50% seeks in typical scenario



Full dataset require 13 seeks
Leaf nodes subject for move **in red**





Second LOD stored in parent nodes
Full dataset is now 10 seeks

Improving throughput

- Data tiles are compressed by pipeline
 - Color tiles are optionally DXT1-compressed
 - Mask compression
 - Source tiles (256x256) are chopped up into smaller pieces (66x66)
 - Empty tile culling
 - Identical (constant value) tile merging
 - Physics materials and destruction depth are packed to four bits and RLE-compressed
 - All tiles are z-compressed

Improving throughput

- A quadtree node has zero *or* four children
- Incomplete quadtree
 - We allow zero *through* four payload children
 - Reduces bundle size by some 20%+



Latency under the rug

- Even with mentioned improvements, streaming is not instant
- General ways of hiding latency
 - CLOD to smooth most streaming LOD transitions
 - Global prioritization helps distribute punishment

Overview


- Scalability – hierarchies, payloads and limitations
- Workflows – realtime in-game editing
- CPU and GPU performance
- Procedural virtual texturing – powerful GPU optimization
- Data streaming – minimizing memory footprint
- **Robustness – global prioritization**
- Procedural mesh generation
- Conclusions

Global prioritization


- Compute priority for each frame and each quadtree node
- Update streaming, virtual texturing and terrain decoration
 - According to priority
- Priority value
 - 1.0: On target
 - One pixel per texel
 - Terrain decoration at specified view distance
 - < 1.0: Node doesn't need payload
 - > 1.0: Node need payload

Priority depends on distance and size

- Closer nodes have higher priority
- Larger nodes have higher priority



Priority modified by culling, updates and speed



Prioritized update algorithm


Low prio

Pool size (keep pool full)

High prio



Nodes



Bidirectional
update

Look for payloads to release

Look for payloads to fetch



Prioritized update cost

- Priority evaluation and sorting done on SPU
 - <1ms
- Update done on PPU
 - <0.5ms

Overview




- Scalability – hierarchies, payloads and limitations
- Workflows – realtime in-game editing
- CPU and GPU performance
- Procedural virtual texturing – powerful GPU optimization
- Data streaming – minimizing memory footprint
- Robustness – global prioritization
- **Procedural mesh generation**
- Conclusions

Mesh

- Mesh generated from heightfield
- Straight-forward tessellation
 - Rendered in patches of 16x16 triangle pairs
 - Projected triangle size roughly constant - support destruction everywhere
 - Blocky crestlines (on console)
- PC, Xenon:
 - Heightfield sampled in vertex shader
- PS3:
 - Vertex shader texture fetch is too slow
 - Heightfield samples stored in vertex attribute
 - Heightfield sampled in pixel shader

Mesh stitching

- A job analyze mesh quadtree and detect LOD switch edges
- Edges are stitched by index permutations
 - Vertices are unchanged



DX11 tessellation

- Displacement mapping from heightfield
 - No additional memory needs (heightfield used as normal map)
- Straightforward hull and domain shaders
- Tessellation factor derived from projected patch edge bounding sphere
 - Tries to maintain a constant screen-space triangle width



Displacement mapping ON

Overview

- Scalability – hierarchies, payloads and limitations
- Workflows – realtime in-game editing
- CPU and GPU performance
- Procedural virtual texturing – powerful GPU optimization
- Data streaming – minimizing memory footprint
- Robustness – global prioritization
- Procedural mesh generation
- **Conclusions**

Conclusions

- Frostbite 2 has a robust and competent terrain system
 - Heightfield, shading, decals, water, terrain decoration
- Most aspects scale well
 - View distance, data resolution, decoration density and distance
- Slick workflow
 - In-game editing
 - Good range of tools
- Good performance (CPU, GPU, memory)
 - Parallelized, streaming, procedural virtual texture

Special thanks

- The Frostbite team
- Black Box
 - Andy Routledge
 - Cody Ritchie
 - Brad Gour
- Criterion
 - Tad Swift
 - Matthew Jones
 - Richard Parr
- Dice
 - Andrew Hamilton

A wide-angle photograph of a majestic mountain range. The foreground is filled with dark, greenish-blue grass and small white flowers. In the middle ground, a valley opens up, leading to a range of mountains with patches of snow and ice clinging to their rocky slopes. The sky above is a mix of bright blue and large, billowing white clouds.

Questions?

mattias.widmark@dice.se
EA Digital Illusions (DICE)